



Lab Session 2 - Final Topics on SQL Using the PostgreSQL DBMS

By Mitchell Welch

University of New England

Reading

- Chapters 6 and 7 from *Fundamentals of Database Systems* by Elmazri and Navathe

Summary

- House Keeping
- String Comparison Operations
- Arithmetic Operations
- Grouping and Aggregate Functions
- Views(Virtual Tables)

Before we Get Started...

- In order to run the examples presented through this practical session, you should import the database you created last week into a new database for this week.
- Create a backup of your `prac_01` database using the `pg_dump` utility:

```
[mwelch8@turing ~]$ pg_dump mwelch8_prac_01 > ~/prac_01.sql
```

- Now create your `prac_03` database for this weeks exercise and import your week 1 database using the `\i` command from within the `psql` client:

```
[mwelch8@turing ~]$ createdb mwelch8_prac_02
```

```
[mwelch8@turing ~]$ psql mwelch8_prac_02 < ~/prac_01.sql
```

```
[mwelch8@turing ~]$ psql mwelch8_prac_02
psql (9.4.4)
Type "help" for help.
```

```
...
```

```
mwelch8_prac_02=> \dt
          List of relations
Schema |      Name      | Type | Owner
-----+-----+-----+-----
public | department     | table | mwelch8
public | dependent      | table | mwelch8
public | dept_locations | table | mwelch8
public | employee       | table | mwelch8
public | project        | table | mwelch8
public | works_on       | table | mwelch8
(6 rows)
```

- Throughout this practical session, you should run each of the example queries and review the results returned from the `COMPANY` database.
- Now you are ready for the prac...

String Comparison Operations

- So far through the practical exercises, we have simply compared string-typed attributes to determine equality.
- SQL and PostgreSQL provide facilities for complex pattern matching string-type attributes.
- Pattern matching* is achieved through the use of the `LIKE` operator.

- The pattern matching facilities can be used for identifying partial matches to patterns.
- The string patterns are specified through the use of two reserved characters:
 - % is used to specify an arbitrary section of 0 or more characters
 - _ is used to specify a single character.
- The reserved characters are simply inserted into the comparison literal on the right side of the LIKE operator.
- Try These:

```
-- Retrieve all employees with an address in Houston,TX
```

```
SELECT fname,lname
FROM employee
WHERE address LIKE '%Houston, TX%'
```

```
-- Find employees that were born in the 1950's
```

```
SELECT fname,lname
FROM employee
WHERE date_part('Year',bdate)::VARCHAR LIKE '__5_'
```

- If you wish to use the '%' and '_' characters as part of a string literal for matching, you will need to escape them in your string using a '\' character:
 - 'AB_CD\%EF' represents the string 'AB_CD%EF'
- A few more interesting examples (source: <http://www.postgresql.org/docs/9.4/static/functions-matching.html>):

```
'abc' LIKE 'abc'      -- true
'abc' LIKE 'a%'       -- true
'abc' LIKE '_b_'      -- true
'abc' LIKE 'c'        -- false
```

- You can mess around with these in postgres:

```
SELECT 'abc' LIKE 'abc' AS answer;
SELECT 'abc' LIKE 'a%' AS answer;
SELECT 'abc' LIKE '_b_' AS answer;
SELECT 'abc' LIKE 'c' AS answer;
```

- PostgreSQL provides POSIX regular expressions through the use of the SIMILAR TO clause.
- SIMILAR TO supports these pattern-matching meta-characters borrowed from POSIX regular expressions:
 - | denotes alternation (either of two alternatives).
 - * denotes repetition of the previous item zero or more times.
 - + denotes repetition of the previous item one or more times.
 - ? denotes repetition of the previous item zero or one time.
 - {m} denotes repetition of the previous item exactly m times.
 - {m,} denotes repetition of the previous item m or more times. {m,n} denotes repetition of the previous item at least m and not more than n times.
 - Parentheses () can be used to group items into a single logical item.
 - A bracket expression [...] specifies a character class, just as in POSIX regular expressions.
- A few Simple ones:

```
'abc' SIMILAR TO 'abc'      -- true
'abc' SIMILAR TO 'a'        -- false
'abc' SIMILAR TO '%(b|d)%'  -- true
'abc' SIMILAR TO '(b|c)%'   -- false
```

- String-type attributes can be concatenated using the || operator:

```
-- Get the full name in a single returned field (with a space between the names)
```

```
SELECT fname||' '||lname AS Full_name
FROM employee;
```

Arithmetic Operations

- The standard arithmetic operators (+, -, *, /) can be applied to numerical values and attributes with numeric domains.

For example:

```
SELECT e.fname, e.lname, 1.1 * e.salary as Increased_salary
FROM employee AS e, works_on AS w, project AS p
WHERE e.ssn = w.essn AND
      w.pno=p.pnumber AND
      p.pname = 'ProductX';
```

- In addition to the basic operators, PostgreSQL provides a full range of builtin functions for performing common operations:

<http://www.postgresql.org/docs/8.1/static/functions-math.html>

Aggregate Functions

- Aggregate functions summarise information from multiple tuples into a single-tuple summary.
- The functions supplied include COUNT, SUM, MAX, MIN and AVERAGE
- A full list of the functions provided in PostgreSQL is available from:
<http://www.postgresql.org/docs/9.1/static/functions-aggregate.html>
- These aggregate functions are applied to the individual attributes specified within a SELECT clause.
- An example for you to try:

```
-- Find the sum of all salaries, the max of all salaries and the min of all salaries.
```

```
SELECT SUM(salary), MAX(salary), MIN(salary)
FROM employee;
```

- We can add a ``WHERE` clause to restrict the records that the aggregate functions are applied to:
- Some examples for you to try:

```
-- Find the sum of all salaries, the max of all salaries and the min of all salaries form the 'Research' department
```

```
SELECT SUM(salary), MAX(salary), MIN(salary)
FROM employee JOIN department on dno=dnumber
WHERE dname='Research';
```

```
-- A Count of all employees
```

```
SELECT COUNT(*) AS "Count of Employees"
FROM employee;
```

```
-- Count the number of distinct salary values in the database
```

```
SELECT COUNT(DISTINCT salary)
FROM employee;
```

Grouping

- In many situations we will need to apply aggregate functions to subgroups of tuples within a relation.
- The GROUP BY operations allows us to partition a relation into a set of non-overlapping groups.
- The GROUP BY operation specifies the attribute(s) to use for creating the subgroups.
 - Tuples with the same value for this specific attribute will be grouped together.
 - Any aggregate functions specified in the query will be carried out across each group. A single tuple will be returned for each of the groups partitioned.

```
-- Retrieve the number of employees and the average salary for each department.
```

```
SELECT dno, COUNT(*), AVG(salary)
FROM employee
GROUP BY dno;
```

- We can restrict rows returned (i.e. groups listed) in the aggregate results through the use of the HAVING clause.
- The HAVING clause puts a condition on the summary information
 - Only groups satisfying the condition will be returned.
- Examples for you to try:

```
-- For each project that has more than 2 employees, list the project number, name and number of employees working on the pro
```

```
SELECT pnumber, pname, COUNT(*)
FROM project, works_on
WHERE pnumber=pno
GROUP BY pnumber, pname
HAVING COUNT(*) > 2;
```

Views(Virtual Tables)

- A *View* is a virtual table derived from the base relations present within the database.
- Views are used to provide an additional layer of abstraction between the implementation of the database and user applications.
 - Changes can be made to the underlying data structure (such as table/column names foreign keys etc.), without changes to the applications
 - Views can simply be updated to return the format expected by the application.
- Views can be created by using the `CREATE VIEW ... AS ...` clause
- Examples for you:

```
-- This view directly inherits the names of the SELECTed attributes from the
-- base tables
```

```
CREATE VIEW works_on1
AS SELECT fname,lname,pname,hours
FROM employee, project, works_on
WHERE ssn=essn AND pno=pnumber;
```

```
-- Look at the contents of your view...
```

```
SELECT * FROM works_on1;
```

```
-- This view renames the attributes and makes use of a GROUP BY clause to
-- return the number of employees and the total salary for each department
```

```
CREATE VIEW dept_info(dept_name, no_of_emps,total_sal)
AS SELECT dname, COUNT(*),SUM(salary)
FROM department, employee
WHERE dnumber = dno
GROUP BY dname;
```

```
-- Look at the contents of your view...
```

```
SELECT * FROM dept_info;
```

- To remove the view, use the `DROP VIEW ...` clause:

```
DROP VIEW dept_info;
```

- The data presented within a view is always up-to-date: If we modify the data in the tables on which a view is based, the data returned in queries to the view will be updated.
- Views data returned by a view can either be retrieved or calculated from the base table on-demand or they can be *Materialised* which involves physically creating a temporary table for the view on the assumption that other queries to the view will follow.
- This is managed within the DBMS.
- Update operations on views are possible if the view directly maps to a base relation.
 - Views that involve aggregate functions can't be updated as there relationship with the base tables can be ambiguous.
 - Views that involve table joins can be ambiguous, as they can map to the underlying tables in multiple ways.
 - Updating a base relation via a view is usually not a good idea as the update may map to multiple update operations on the underlying base relations.
 - These update operations can be unpredictable.

Exercises For You

Question 1

Specify the following queries on the COMPANY database in SQL. Execute your query and review the results.

- For each department whose average employee salary is more than \$30,000, retrieve the department name and the number of employees working for that department.
- Retrieve all employees with a last name that starts the the character 'S'.
- Find all employees that work on a project with the the word 'Product' in its name.

Question 2

Specify the following views in SQL on the COMPANY database. Execute your query and review the results.

- A view that has the department name, manager name, and manager salary for every department.

b) A view that has the employee name, supervisor name, and employee salary for each employee who works in the 'Research' department.

c) A view that has the project name, controlling department name, number of employees, and total hours worked per week on the project for each project.

Exercise 3

Using appropriate SQL DDL statements, create an SQL script (i.e. a file with the .sql extension) to implement that following schema:

STUDENT

| Name | Student_number | Class | Major |
|-------|----------------|-------|-------|
| Smith | 17 | 1 | CS |
| Brown | 8 | 2 | CS |

COURSE

| Course_name | Course_number | Credit_hours | Department |
|---------------------------|---------------|--------------|------------|
| Intro to Computer Science | CS1310 | 4 | CS |
| Data Structures | CS3320 | 4 | CS |
| Discrete Mathematics | MATH2410 | 3 | MATH |
| Database | CS3380 | 3 | CS |

SECTION

| Section_identifier | Course_number | Semester | Year | Instructor |
|--------------------|---------------|----------|------|------------|
| 85 | MATH2410 | Fall | 07 | King |
| 92 | CS1310 | Fall | 07 | Anderson |
| 102 | CS3320 | Spring | 08 | Knuth |
| 112 | MATH2410 | Fall | 08 | Chang |
| 119 | CS1310 | Fall | 08 | Anderson |
| 135 | CS3380 | Fall | 08 | Stone |

GRADE_REPORT

| Student_number | Section_identifier | Grade |
|----------------|--------------------|-------|
| 17 | 112 | B |
| 17 | 119 | C |
| 8 | 85 | A |
| 8 | 92 | A |
| 8 | 102 | B |
| 8 | 135 | A |

PREREQUISITE

| Course_number | Prerequisite_number |
|---------------|---------------------|
| CS3380 | CS3320 |
| CS3380 | MATH2410 |
| CS3320 | CS1310 |

Figure 1.2

A database that stores student and course information.

Identify and implement the primary and foreign key values from the data supplied.

Finish your script by implementing a series of `INSERT` statements to load this data into the database. Run your script within your `prac_02` database using the `psql` client's `\i` command.